

CS 251 / EE 255 REAL-TIME EMBEDDED SYSTEMS
HW Project #2: Task Scheduling and Monitoring
Group16

Group members

Hengshuo Zhang
hzhan402@ucr.edu

Zhaoze Sun
zsun114@ucr.edu

Writeup

1. Problem 1

Steps for set up:

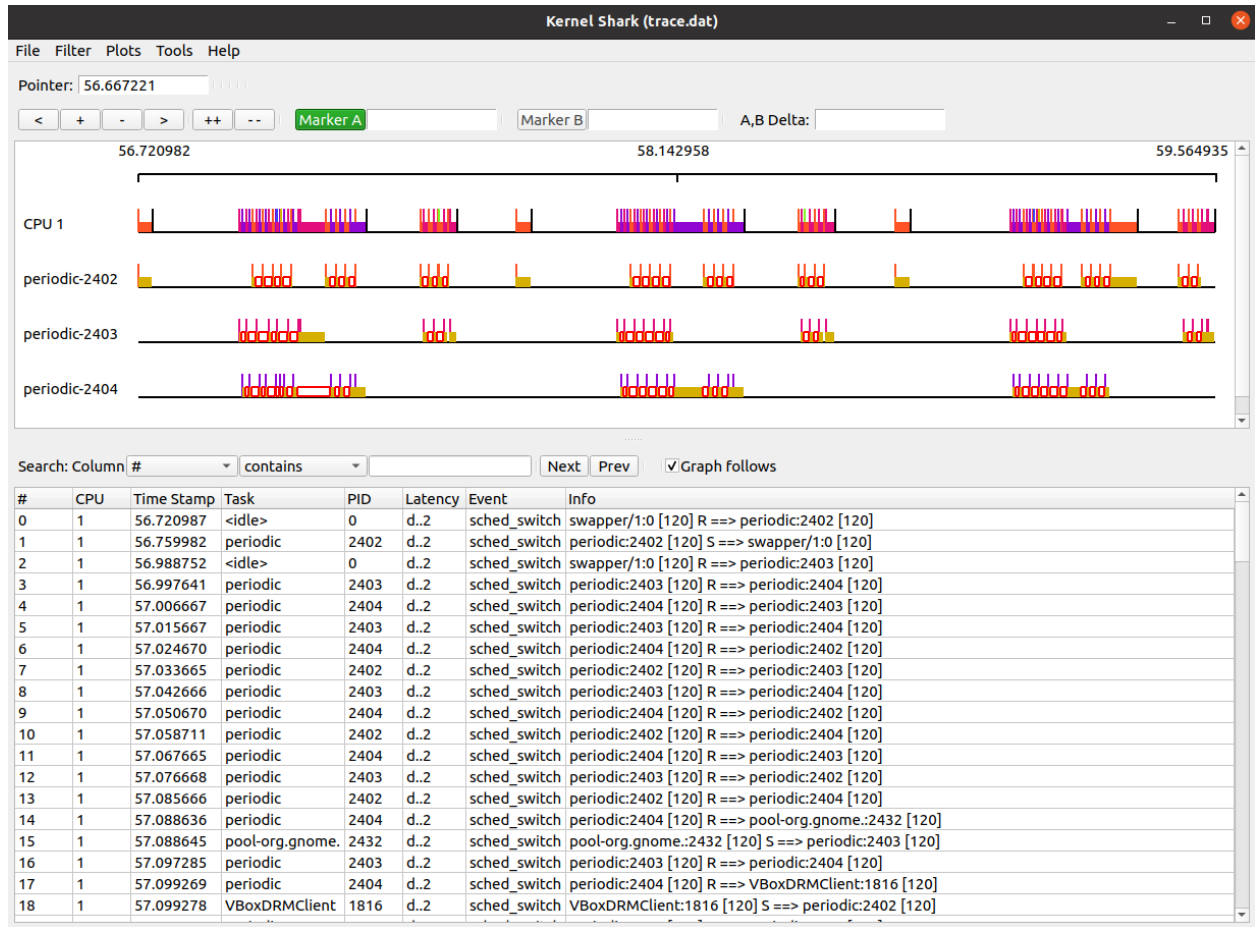
```
rex@rex:~/rtes/proj2/apps/periodic$ ./periodic 40 250 1 & ./periodic 60 500 1 & ./periodic 200 1000 1 &  
[1] 2402  
[2] 2403  
[3] 2404  
rex@rex:~/rtes/proj2/apps/periodic$ PID: 2402, C: 40, T: 250, CPUID: 1  
PID: 2403, C: 60, T: 500, CPUID: 1  
PID: 2404, C: 200, T: 1000, CPUID: 1
```

```
// Run three instances of the periodic program  
> ./periodic 40 250 1 & ./periodic 60 500 1 & ./periodic 200  
1000 1 &
```

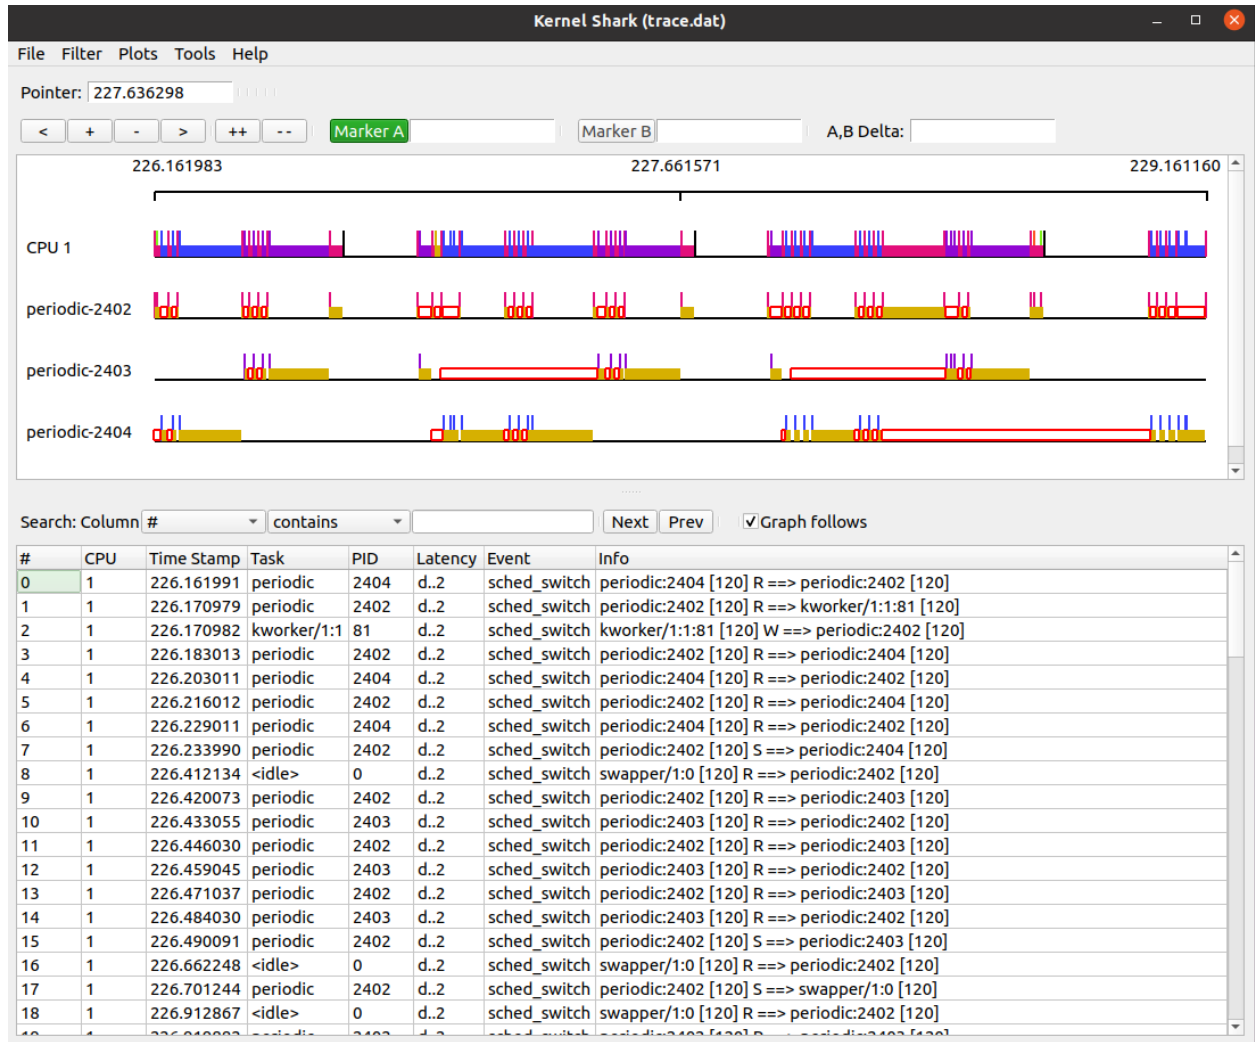
The first instance has PID 2402, with a period of 40ms, a computation time of 250ms.

The second instance has PID 2403, with a period of 60ms, a computation time of 500ms.

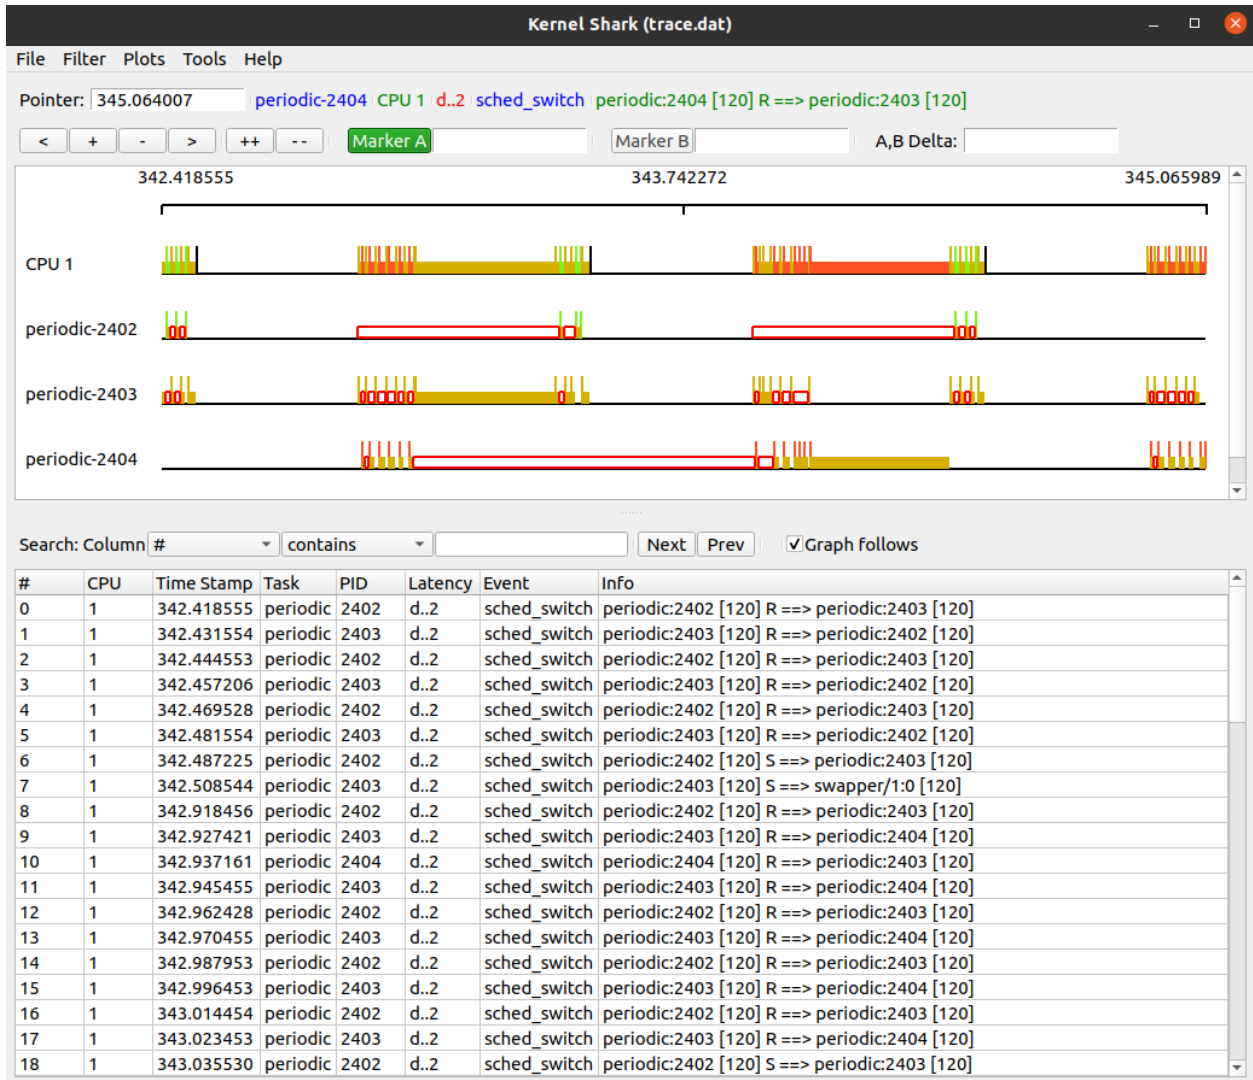
The third instance has PID 2404, with a period of 200ms, a computation time of 1000ms.



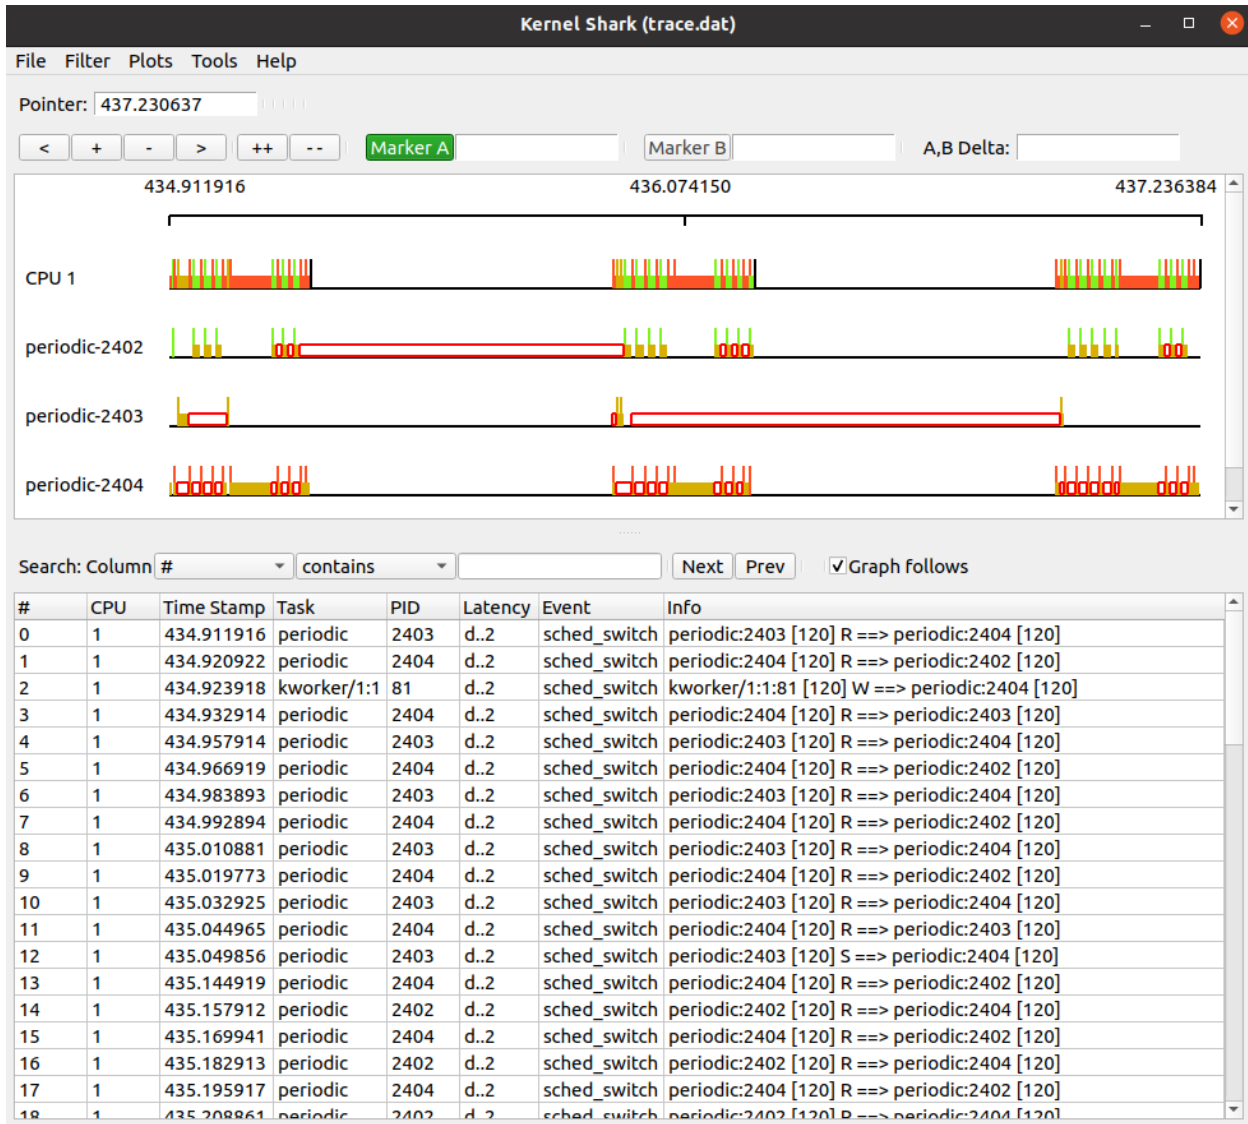
```
// record all tasks
> sudo trace-cmd record -e sched_switch -P 2402 -P 2403 -P 2404
sleep 3
// displays all events recorded
> kernelshark
```



```
// Record scheduling events of 2402 for 3s
> sudo trace-cmd record -e sched_switch -P 2402 sleep 3
// displays all events recorded
> kernelshark
```



```
// Record scheduling events of 2403 for 3s
> sudo trace-cmd record -e sched_switch -P 2403 sleep 3
// displays all events recorded
> kernelshark
```



```
// Record scheduling events of 2404 for 3s
> sudo trace-cmd record -e sched_switch -P 2404 sleep 3
// displays all events recorded
> kernelshark
```

2. Problem 2

Steps for set up:

```
> sudo chrt -f -p 80 2402  
> sudo chrt -f -p 79 2403  
> sudo chrt -f -p 78 2404
```

High Priority: 2402 (80)

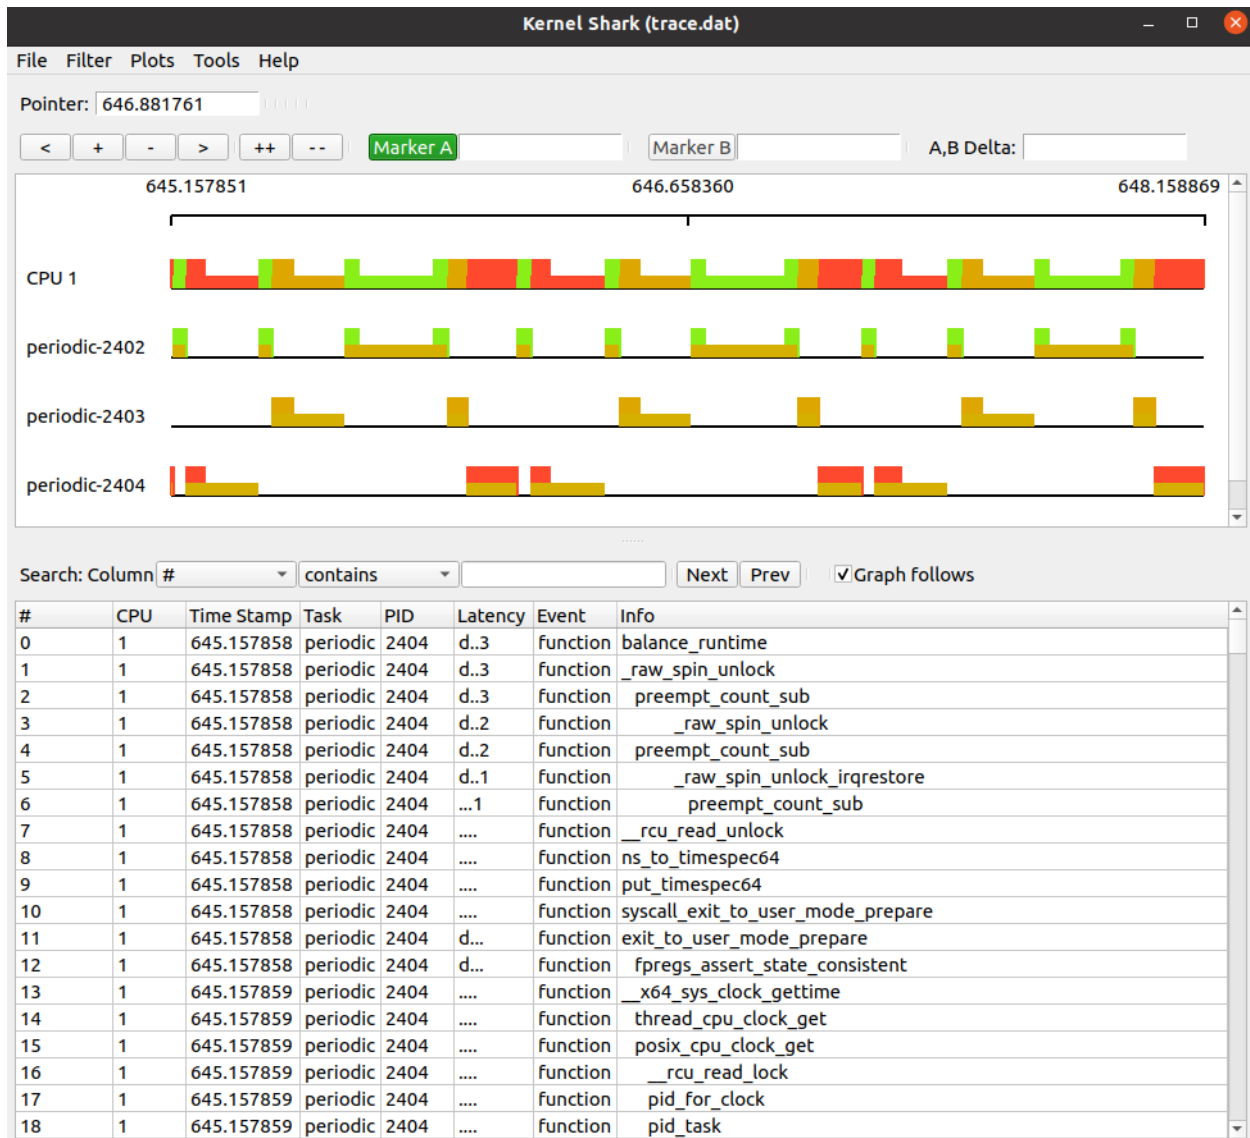
Set the real-time priority of 2402 to 80 and implement the SCHED_FIFO scheduling policy.

Mid Priority: 2403 (79)

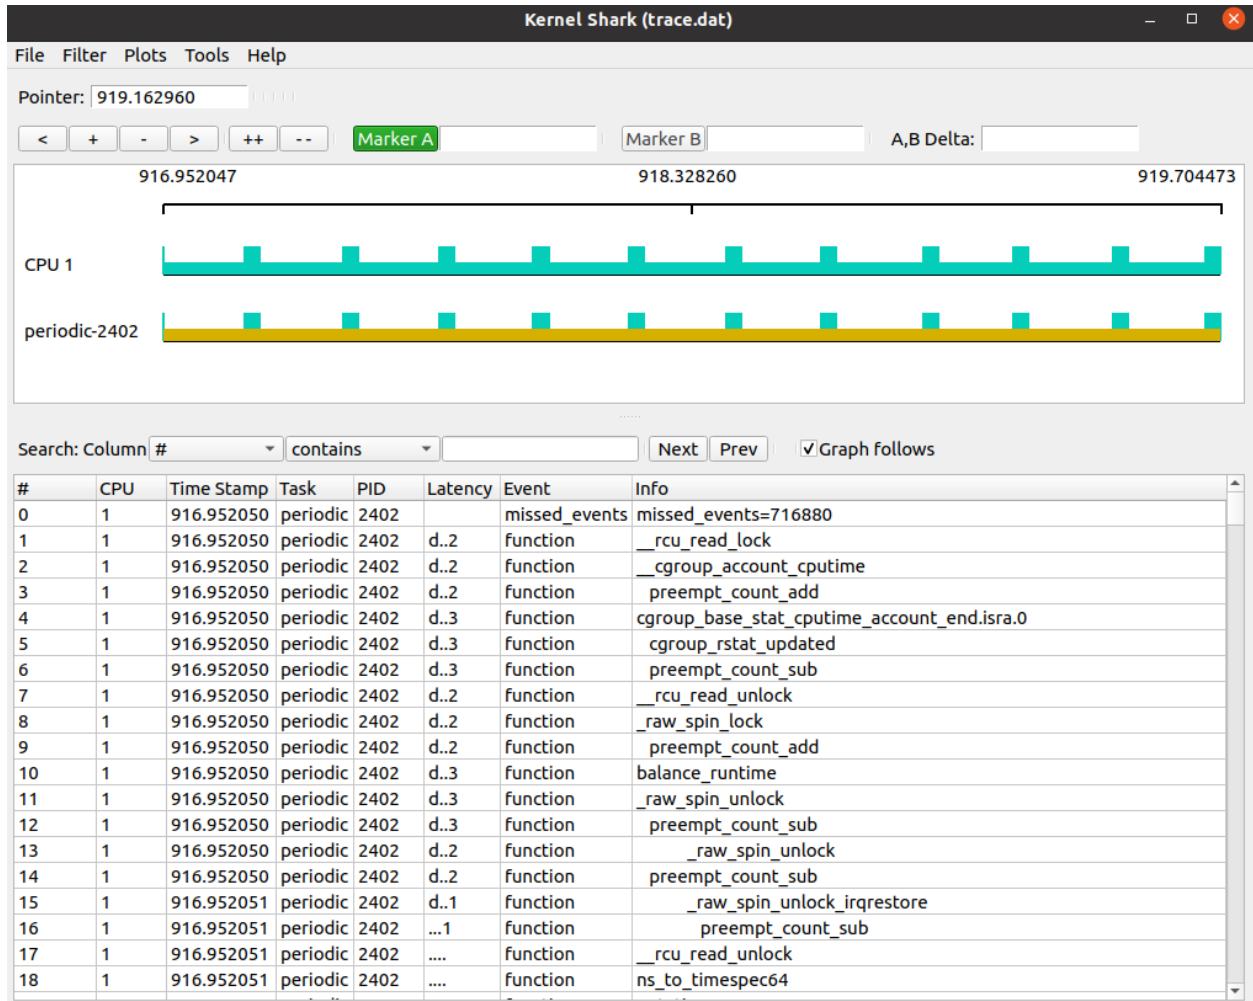
Set the real-time priority of 2403 to 79 and implement the SCHED_FIFO scheduling policy.

Low Priority: 2404 (78)

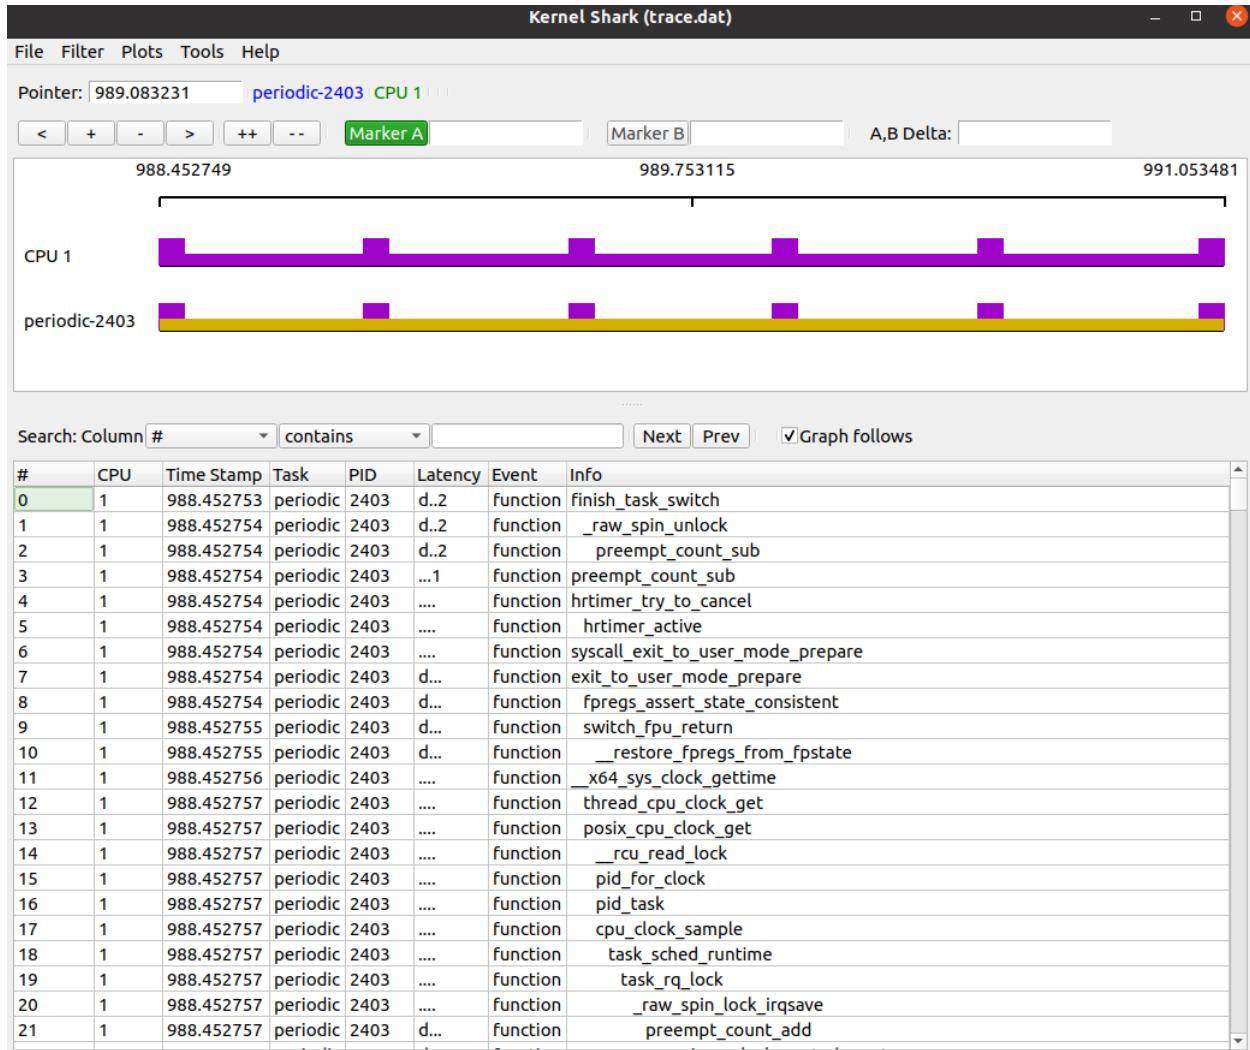
Set the real-time priority of 2404 to 78 and implement the SCHED_FIFO scheduling policy.



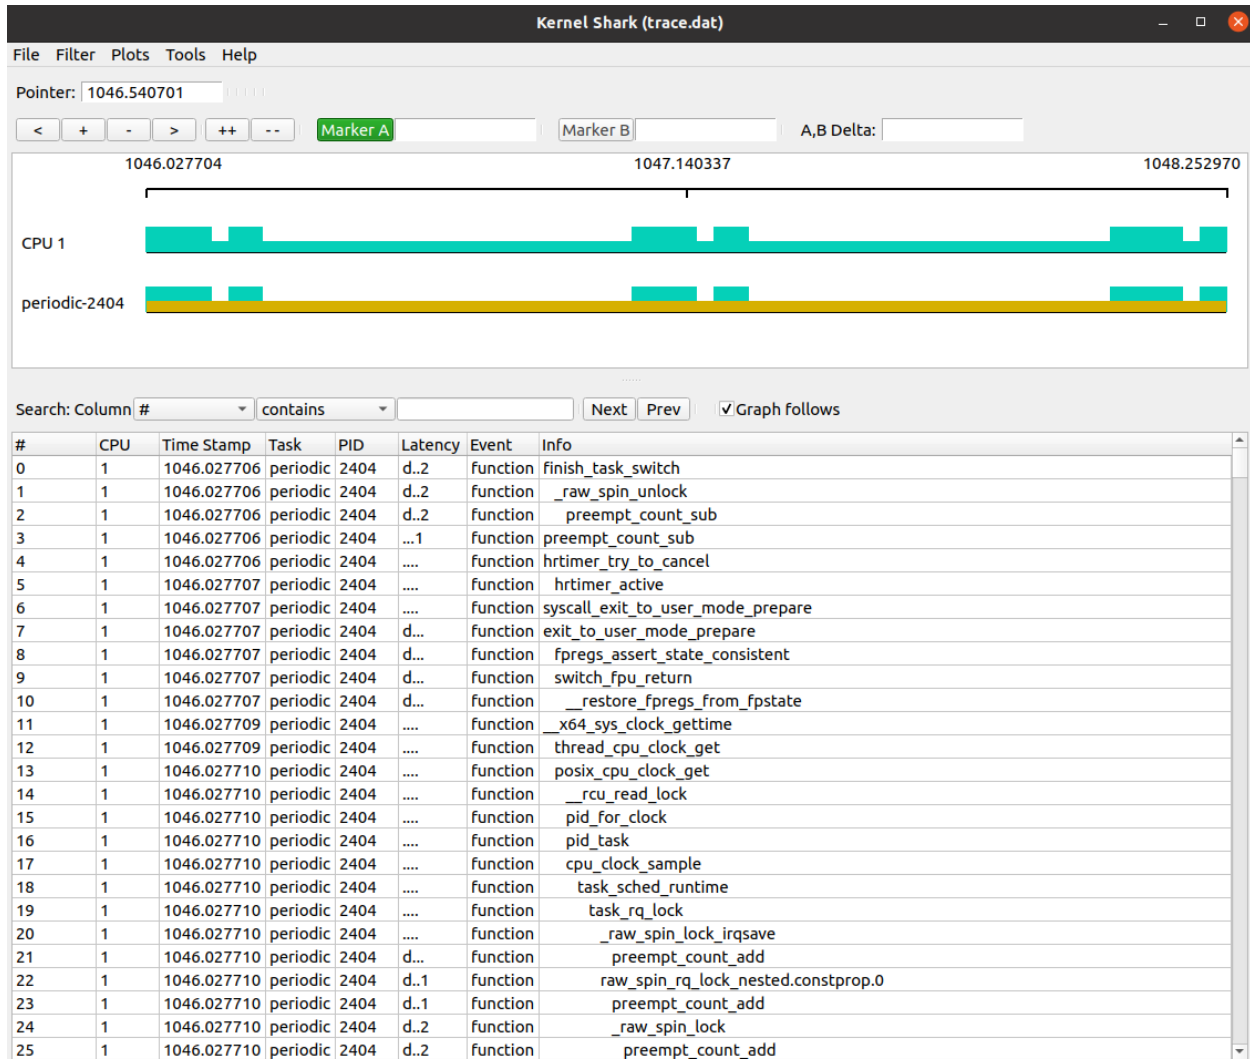
```
> sudo trace-cmd record -p function -P 2402 -P 2403 -P 2404
sleep 3
> kernelshark
```



```
> sudo trace-cmd record -p function -P 2402 sleep 3
> kernelshark
```

```
> sudo trace-cmd record -p function -P 2403 sleep 3
> kernelshark
```



```
> sudo trace-cmd record -p function -P 2404 sleep 3
> kernelshark
```


3. Problem 3

In Task 1, each task is executed using the regular scheduling policy. The execution trace will show that each task is executed at a given time period with a delay equal to or less than the system jitter.


The periodic program creates a thread for each instance and uses the `clock_nanosleep` system call to put the thread to sleep for a period of time. When the thread wakes up, it executes the task, which is a simple busy-wait cycle that consumes CPU time for a specified duration C . After the task completes, the thread returns to the sleep state for the remainder of the time until the next cycle T .

In Task 2, each task is assigned a real-time priority via `SCHED_FIFO`. The execution trace will show that the higher priority task is executed before the lower priority task. Tasks with lower priority may have delays beyond their cycle time because the system always prioritizes tasks with higher priority. However, tasks with higher priority may face longer blocking times from the operating system or other tasks.

Real-time priority is a value that determines the order in which tasks are executed by the kernel scheduler. A higher priority value indicates a more critical task that needs to be executed before a lower priority task. With the `SCHED_FIFO` policy, this means that tasks with higher priority will run to completion before any lower priority tasks run. The priority values range from 1 to 99, where 99 is the highest priority.



Regular Scheduling



SCHED_FIFO

In a regular scheduling diagram, there is no real-time priority assignment in the cycle, and we are not able to find exactly the same two execution traces. For the analysis of the execution traces, we are not able to visualize whether the task has completed in the cycle or has been delayed or lost.

For SCHED_FIFO, the periodicity of execution traces will be scheduled with real-time priority. We can intuitively understand the execution of tasks under the same CPU by a single cycle, and with the influence of priority, we will have a more intuitive understanding of the priority relationship between tasks.

Here is summary of using different strategies:

Method	Pros	Cons
Regular Scheduling Policy	Simple and easy to implement	May not be suitable for real-time systems or applications that require timely responses to external events
	Tasks execute in a predictable and deterministic manner	Tasks may experience jitter or delay if the system is overloaded or if other high-priority tasks are executing at the same time
		Cannot prioritize critical tasks over less important tasks
Real-Time Priority Assignments with SCHED_FIFO	Provides a way to prioritize critical tasks over less important tasks	Requires knowledge and understanding of the kernel scheduler and real-time priorities
	Enables real-time systems or applications to respond quickly to external events	Assigning high priorities to tasks can lead to priority inversion, where a low-priority task blocks a higher-priority task
	Tasks are executed in a predictable and deterministic manner	Assigning high priorities to tasks can also lead to starvation, where a low-priority task is not executed at all

4. Summary of The Contributions

Overall:

Hengshuo Zhang 50%

Zhaoze Sun 50%

Breakdown to each Assignment

4.2	Setting Task's Timing Parameters	Hengshuo/Zhaoze
4.3	Printing Task's Timing Parameters	Hengshuo/Zhaoze
4.4	Admission Control	Hengshuo/Zhaoze
4.5	Computation Time Tracking	Hengshuo/Zhaoze
4.6	Periodic Task Support	Hengshuo/Zhaoze
4.8	Writeup	Hengshuo/Zhaoze

Readme:

The (set_rtmon/print_rtmon/cancel_rtmon) system calls in 4.2 to 4.5 are saved in the directory path (... rtes/proj2/kernel/rtmon.c).

For testing purposes, the test.c file in the apps folder will be available for testing the three system calls. The path to the test file is (... rtes/proj2/apps/periodic/testfile/test.c).

Demo (set_rtmon/print_rtmon/cancel_rtmon) system calls:

```
> gcc -o test test.c
> sudo ./test 200 1200
> dmesg
```

```
rex@rex:~/rtes/proj2/apps/periodic/testfile$ gcc -o test test.c
rex@rex:~/rtes/proj2/apps/periodic/testfile$ sudo ./test 200 1200
[sudo] password for rex:
Setting up monitoring for PID 2484, C = 200 ms, T = 1200 ms...
Monitoring set up successfully. Wait for 3 seconds to print and cancel usage information.
Monitoring print/cancel successfully.
rex@rex:~/rtes/proj2/apps/periodic/testfile$ dmesg | tail -1
[ 90.427797] print_rtmon: PID 2484, C 200 ms, T 1200 ms, usage 581 ms
```

We used a separate .c file for the system calls (wait_until_next_period) in 4.6, which will show the logic more clearly and reduce pointer problems at runtime. The file path is (... rtes/proj2/kernel).

Based on the test structure mentioned in 4.6, we wrote a separate test file to test the system call. The path of the test file is (... rtes/proj2/apps/periodic/testfile/demo.c).

Demo (wait_until_next_period) system calls:

```
> gcc -o demo demo.c
```

```
> sudo ./demo
```

```
rex@rex:~/rtes/proj2/apps/periodic/testfile$ gcc -o demo demo.c
rex@rex:~/rtes/proj2/apps/periodic/testfile$ sudo ./demo
[sudo] password for rex:
Setting RT mode...
RT mode set!
Do computation... (loop 0)
Do computation... (loop 1)
Do computation... (loop 2)
Do computation... (loop 3)
Do computation... (loop 4)
^C
```